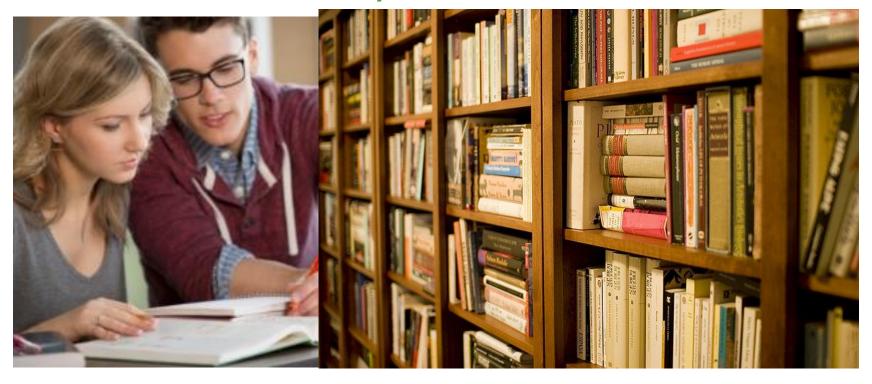
## CHAPTER 7-2

Memory

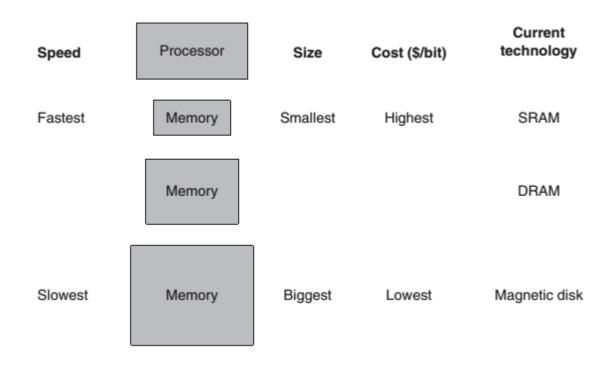
By Pattama Longani Collage of arts, media and Technology

### Books in a library



 $http://www.gettyimages.com/detail/photo/college-students-studying-at-table-royalty-free-image/485207369 \\ http://commons.wikimedia.org/wiki/File:Bookshelf.jpg$ 

- •A memory hierarchy consists of multiple levels of memory with different speeds and sizes.
- •The faster memories are more expensive per bit than the slower memories and thus are smaller.
- •performance is the major reason for having a memory hierarchy, the time to service hits and misses is important.

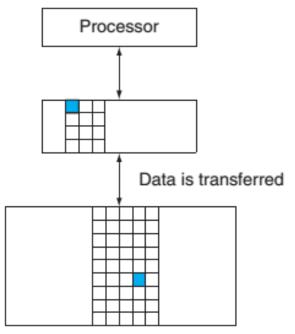


 data is copied between only two adjacent levels at a time

- •block (or line) The minimum unit of information
- •hit rate The fraction of memory accesses found in a level of the memory hierarchy.
- •miss rate The fraction of memory accesses not found in a level of the memory hierarchy.

#### •Hit time is the time to access the upper level of the memory hierarchy, which includes the time needed to determine whether the access is a hit or a miss

 miss penalty is the time to replace a block in the upper level with the corresponding block from the lower level, plus the time to deliver this block to the processor



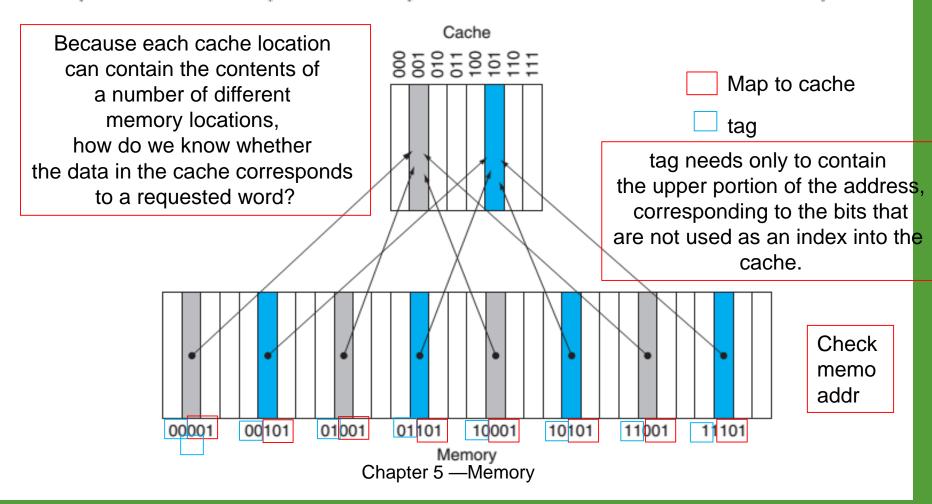
•Because the upper level is smaller and built using faster memory parts, the hit time will be much smaller than the time to access the next level in the hierarchy,

### memory systems affection

- how the operating system manages memory and I/O
- how compilers generate code
- how applications use the computer

#### Cache

(Block address) modulo (Number of blocks in the cache)



### Accessing a Cache

Index field used is used as an address to reference the cache

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	N		
111	N		

a. The initial state of the cache after power-on

Index	V	Tag	Data
000	N		
001	N		
010	Υ	11 <sub>two</sub>	Memory (11010 <sub>two</sub> )
011	N		
100	N		
101	N		
110	Υ	10 <sub>two</sub>	Memory (10110 <sub>two</sub> )
111	N		

c. After handling a miss of address (11010<sub>two</sub>)

Index	V	Tag	Data
000	Υ	10 <sub>two</sub>	Memory (10000 <sub>two</sub> )
001	N		
010	Υ	11 <sub>two</sub>	Memory (11010 <sub>two</sub> )
011	Υ	00 <sub>two</sub>	Memory (00011 <sub>two</sub> )
100	N		
101	N		
110	Υ	10 <sub>two</sub>	Memory (10110 <sub>two</sub> )
111	N		

e. After handling a miss of address (00011two)

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	Y	10 <sub>two</sub>	Memory (10110 <sub>two</sub> )
111	N		

b. After handling a miss of address (10110<sub>two</sub>)

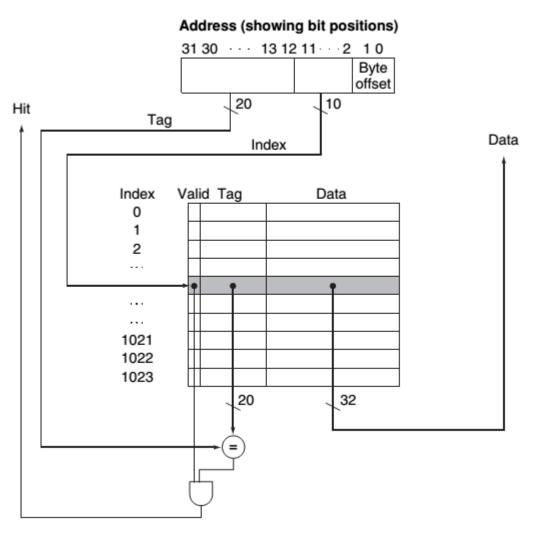
Index	V	Tag	Data
000	Y	10 <sub>two</sub>	Memory (10000 <sub>two</sub> )
001	N		
010	Y	11 <sub>two</sub>	Memory (11010 <sub>two</sub> )
011	N		
100	N		
101	N		
110	Y	10 <sub>two</sub>	Memory (10110 <sub>two</sub> )
111	N		

d. After handling a miss of address (10000<sub>two</sub>)

Index	V	Tag	Data
000	Y	10 <sub>two</sub>	Memory (10000 <sub>two</sub> )
001	N		
010	Y	10 <sub>two</sub>	Memory (10010 <sub>two</sub> )
011	Y	OO <sub>two</sub>	Memory (00011 <sub>two</sub> )
100	N		
101	N		
110	Y	10 <sub>two</sub>	Memory (10110 <sub>two</sub> )
111	N		

f. After handling a miss of address ( $10010_{two}$ )

### **Accessing a Cache**



#### MIPS cache

- •Since words are aligned to multiples of four bytes, the least significant two bits of every address specify a byte within a word.
- •Hence, the **least significant two bits are ignored** when selecting a word in the block.
- •The total number of bits needed for a cache is a function of the **cache size** and the **address size**.
- •Cache includes both the storage for the data and the tags.

#### MIPS cache

- 32-bit byte addresses
- A direct-mapped cache
- $\blacksquare$  The cache size is  $2^n$  blocks, so *n* bits are used for the index
- The block size is  $2^m$  words  $(2^{m+2}$  bytes), so m bits are used for the word within the block, and two bits are used for the byte part of the address

the size of the tag field is

Byte offset 
$$32 - (n + m + 2)$$
.

The total number of bits in a direct-mapped cache is

 $2^n \times$  (block size + tag size + valid field size).

### Ex: Bits in a Cache

How many total bits are required for a direct-mapped cache with 16 KB of data and 4-word blocks, assuming a 32-bit address?

- 16 KB = 4K ( $2^{12}$ ) words. Different data
- With a block size of 4 words ( $2^2$ ), there are 1024 ( $2^{10}$ ) blocks. > cache size n = 10
- Each block has 4 \*32 or 128 bits of data
- Tag = 32 10 2 2 bits
- valid bit = 1 bit.
- Thus, the total cache size is

$$2^{10}$$
 \*  $(4$  \*32 +  $(32 - 10 - 2 - 2) + 1) =  $2^{10}$  *147 = 147 **Kbits**$ 

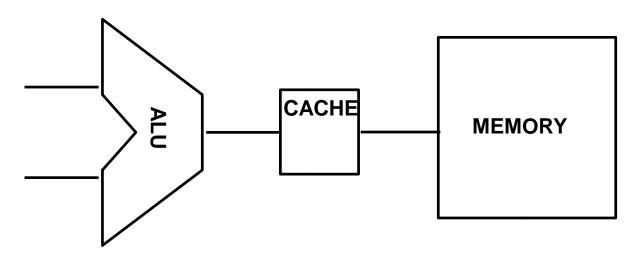
- •The cache miss handling is done in collaboration with the processor control unit and with a separate controller that initiates the memory access and refills the cache.
- •The processing of a cache miss creates a pipeline stall.
- •As opposed to an **interrupt**, which would require saving the state of all registers.

- •For a cache miss, we can stall the entire processor, essentially freezing the contents of the temporary and programmer-visible registers, while we wait for memory.
- •More sophisticated out-of-order processors can allow execution of instructions while waiting for a cache miss, but we'll assume inorder processors that stall on cache misses in this section.

Steps to be taken on an instruction cache miss:

- •1. Send the original PC value (PC= PC-4) to the memory.
- •2. Instruct main memory to perform a read and wait for the memory to complete its access.

- •3. Write the cache entry,
  - putting the data from memory in the data portion
  - writing the upper bits of the address into the tag field
  - turning the valid bit on
- •4. Restart the instruction execution at the first step
  - refetch the instruction, this time finding it in the cache.



•write-through A scheme in which writes always update both the cache and the next lower level of the memory hierarchy, ensuring that data is always consistent between the two.

•Write-through (every write causes the data to be written to main memory) can takes a long time.

#### For example

- •suppose 10% of the instructions are stores. If the CPI without cache misses was 1.0, spending 100 extra cycles on every write would lead to a CPI of 1.0 + 100 \*10% = 11
- •Reducing performance by more than a factor of 10.

- One solution to this problem is to use a write buffer.
- •A write buffer stores the data while it is waiting to be written to memory.
  - •After writing the data into the cache and into the write buffer, the processor can continue execution.
  - •When a write to main memory completes, the entry in the write buffer is freed.
  - •If the write buffer is full when the processor reaches a write, the processor must stall until there is an empty position in the buffer

- •if the rate at which the memory can complete writes is less than the rate at which the processor is generating writes, no amount
- •Stalls may still occur if the rate at which writes are generated may be *less* than the rate at which the memory can accept them.
- •To reduce the occurrence of such stalls, processors usually increase the depth of the write buffer beyond a single entry.

- •Another alternative to a write-through scheme is a scheme called write-back
  - •Handles writes by updating values only to the block in the cache, then writing the modified block to the lower level of the hierarchy when the block is replaced.
  - •Hardware modifier
  - more complex

- •reducing the miss rate by reducing the probability that two different memory blocks will contend for the same cache location.
- •reduces the miss penalty by adding an additional level to the hierarchy, multilevel caching technique

•CPU time can be divided into the clock cycles that the CPU spends executing the program and the clock cycles that the CPU spends waiting for the memory system.

CPU time = (CPU execution clock cycles + Memory-stall clock cycles)
× Clock cycle time

- •We make an assumption that the memory-stall clock cycles come primarily from cache misses.
- Memory-stall clock cycles can be defined as the sum of the stall cycles coming from reads plus those coming from writes:

Memory-stall clock cycles = Read-stall cycles + Write-stall cycles

$$Read-stall\ cycles = \frac{Reads}{Program} \times Read\ miss\ rate \times Read\ miss\ penalty$$

$$Write-stall\ cycles = \left(\frac{Writes}{Program} \times Write\ miss\ rate \times Write\ miss\ penalty\right) \\ + Write\ buffer\ stalls$$

- •Because the write buffer stalls depend on the proximity of writes, and not just the frequency, it is not possible to give a simple equation to compute such stalls.
- •the write buffer stalls will be small, and we can safely **ignore them**.

•In most write-through cache organizations, the read and write miss penalties are the same, we can combine the reads and writes by using a single miss rate and the miss penalty

$$Memory-stall\ clock\ cycles = \frac{Memory\ accesses}{Program} \times Miss\ rate \times Miss\ penalty$$

$$Memory-stall\ clock\ cycles = \frac{Instructions}{Program} \times \frac{Misses}{Instruction} \times Miss\ penalty$$

 Assume the miss rate of an instruction cache is 2% and the miss rate of the data cache is 4%. If a processor has a CPI of 2 without any memory stalls and the miss penalty is 100 cycles for all misses, determine how much faster a processor would run with a perfect cache that never missed. Assume the frequency of all loads and stores is 36%.

Instruction miss cycles =  $I \times 2\% \times 100 = 2.00 \times I$ 

Data miss cycles =  $I \times 36\% \times 4\% \times 100 = 1.44 \times I$ 

The total number of memory-stall cycles is 2.00 I + 1.44 I = 3.44 I

Total CPI = 
$$2 + 3.44 = 5.44$$

$$\frac{\text{CPU time with stalls}}{\text{CPU time with perfect cache}} = \frac{I \times \text{CPI}_{\text{stall}} \times \text{Clock cycle}}{I \times \text{CPI}_{\text{perfect}} \times \text{Clock cycle}}$$
$$= \frac{\text{CPI}_{\text{stall}}}{\text{CPI}_{\text{perfect}}} = \frac{5.44}{2}$$

The performance with the perfect cache is better by  $\frac{5.44}{2} = 2.72$ .

 Memory system is related to many part of computer.

What happens if the processor is made faster, but the memory system is not?

Ex: Suppose we speed-up the computer in the previous example by reducing its CPI from 2 to 1 without changing the clock rate, which might be done with an improved pipeline.

•The system with cache misses would then have a CPI of 1 + 3.44 = 4.44

•Speed up = 
$$\frac{4.44}{1} = 4.44$$

The amount of execution time spent on memory stalls would have risen from

to

$$3.44/4.44 = 77\%$$
.

- •To capture the fact that the time to access data for both hits and misses affects performance, designers sometime use average memory access time (AMAT) as a way to examine alternative cache designs.
- •Average memory access time is the average time to access memory considering both hits and misses and the frequency of different accesses; it is equal to the following:

## Example

•Find the AMAT for a processor with a 1 ns clock cycle time, a miss penalty of 20 clock cycles, a miss rate of 0.05 misses per instruction, and a cache access time (including hit detection) of 1 clock cycle. Assume that the read and write miss penalties are the same and ignore other write stalls

AMAT = Time for a hit + Miss rate  $\times$  Miss penalty = 1 + 0.05  $\times$  20 = 2 clock cycles or 2 ns.